

# Object Detection with Jetson TX1: Challenges and Insights

Oron Port      Gil Shomron  
{soronpo, gilsho}@tx.technion.ac.il  
Electrical Engineering  
Technion — Israel Institute of Technology

## ABSTRACT

In this short paper we summarize

## 1. INTRODUCTION

Object detection is a major rule in many different devices, from mobile phones, cameras, and IoT devices, to drones, and autonomous cars. With the aid of machine learning, detecting an object also involves its classification. Fast detection, and accurate classification... I don't know.

GPU is built for throughput. Implemented with a large number of cores (*streaming multiprocessors* (SM)) its parallelism enables algorithms, such as object detection and machines learning, to run faster than on a CPU.

With increasing demand for low-energy modules, NVIDIA started manufacturing an embedded platform called Jetson. We have received the Jetson TX1 to explore the performance of different object detection algorithms.

In this paper we will discuss and analyze the performance of two algorithms: (1) *Single Shot MultiBox Detector* (SSD) [1], as implemented with Caffe, and (2) *You Only Look Once* (YOLO) [2].

## 2. RESULTS

To compare Caffe SSD performance versus YOLO performance on Jetson TX1, we used their supplied applications for object detection on image files. We downloaded 600 images from MSCOCO dataset [3], and measured the total execution time. We made three comparison between the two algorithms: when both run on (1) CPU, (2) GPU without cuDNN, and (3) GPU with cuDNN<sup>1</sup>.

Figure 1 shows the execution time of both SSD and YOLO for each of the options described above. Not surprisingly, running object detection algorithms on a CPU is not as efficient as running them on a GPU. The

<sup>1</sup>The NVIDIA CUDA Deep Neural Network library (cuDNN) is a GPU-accelerated library of primitives for deep neural networks. cuDNN provides highly tuned implementations for standard routines such as forward and backward convolution, pooling, normalization, and activation layers.

GPU built-in parallelism is advantageous for such algorithms. [TODO: comparison between SSD and YOLO on CPU]

[TODO: comparison between GPU with and without cuDNN]

[TODO: comparison between SSD and YOLO on GPU]

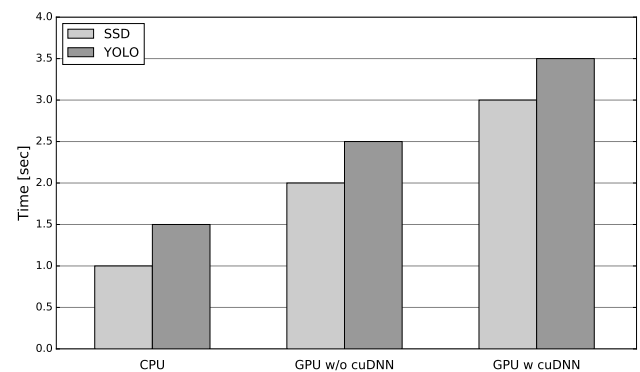


Figure 1: Execution time comparison

## 3. PROFILING

Figure 2 shows the profiling output of object detection of a single image, using SSD and YOLO. Profiling was done using Nvidia Visual Profiler.

One of Jetson's strengths is that its CPU and GPU share the same memory. Therefore, host-to-device (HtoD) and device-to-host (DtoH) copies can be optimized to zero-copy [4]. Zero-copy will have no effect on SSD and YOLO, since memory copies are not frequent, therefore, there is no speedup potential. In addition, there are no overlapping opportunities between HtoD, DtoH, and the compute kernels, since there are almost no memory copies.

In terms of compute intensity, both algorithms use 75% of their compute time in cuDNN kernels. There is no parallelism between kernels. SSD and YOLO fetch and analyze each image in serial, so theoretically, one can analyze a couple of images in parallel. However, Jetson TX1 has only 2 SMs, with 128 CUDA cores each,

therefore, adding additional threads to the system, or running two or more object detection processes in parallel, achieves no performance gains.

## 4. INSTALLATION AND EXECUTION

Some aspects of the Jetson installation process are not trivial. In this section, we share some installation tips, gathered during this project. It is our hope that these tips will others to avoid our obstacles.

### 4.1 Jetson TX1 Hardware Kit

The kit includes most of the hardware required, but is **missing** the following:

- Power Cable – Must
- Network Cable – Must
- External USB Web Camera – Must (if camera is required for the application).
- USB Keyboard – Optional
- USB Mouse – Optional
- HDMI Cable – Optional (with an HDMI supported monitor)
- USB Flash Drive – Optional (If extra space is required)

The on-board camera is not supported well, while a simple USB web camera connected to the USB worked right out of the box.

The optional components are required to debug the system. All the work can be done via SSH, but our initial system bring-up had unexpected problems. Hopefully, using the tips we layout here, others will be able to avoid this in the future.

### 4.2 Jetpack SDK Installation

Download the latest Jetpack SDK [5]. nVidia Developer membership is required (registration is free). It is best to connect the Jetson to a router and SSH to it via internal network. The initial setup requires a host computer connected via USB to the Jetson. The host **must** be an Ubuntu 14.04 machine. It is possible (and recommended) to use a Virtual Machine software with the proper image. We used Virtual Box [6] with the Ubuntu 14.04.05 Trusty image [7]. Do not confuse with the 64-bit Ubuntu 16.04 version, which is installed on the Jetson.

The Jetpack installs an optimized OpenCV library, but of an old version. We attempted to update to a newer version, but had too many complication. It is our recommendation to leave the installed version as is.

The Jetson has very little installation space ( 14GB). It is possible to remove the installation files to free space after the installation. In any case, it is recommended to have an EXT3/EXT4 formatted USB flash drive available (FAT/NTFS formatted drives do not work as well).

### 4.3 Caffe Installation

The SSD Caffe[8] is a fork of the Caffe library[9]. When cloning/downloading the project it is required to use the *ssd* branch.

During Caffe compilation we suffered from a lot system restarts. We figured out that the CPU fan is turned-on by default, which cause the system to heat-up. **The fan must be turned on!**

**Makefile.config.** To use cuDNN acceleration:

```
USE_CUDNN := 1
```

Tegra X1 has CUDA capability 5.3, therefore append to *CUDA\_ARCH*:

```
CUDA_ARCH := -gencode arch=compute_53 ,  
↳ code=sm_53
```

HDF5 directories should be added to *INCLUDE\_DIRS* and *LIBRARY\_DIRS*:

```
INCLUDE_DIRS := $(PYTHON_INCLUDE) /usr/  
↳ local/include /usr/include/hdf5/  
↳ serial  
LIBRARY_DIRS := $(PYTHON_LIB) /usr/  
↳ local/lib /usr/lib/aarch64-linux-  
↳ gnu/serial
```

**Makefile.** HDF5 libraries need to be added to the Makefile also:

```
LIBRARIES += glog gflags protobuf  
↳ boost_system boost_filesystem m  
↳ hdf5_serial_hl hdf5_serial
```

**Python.** Caffe also has Python libraries. Running Python scripts can fail due to unset Python path. By running:

```
export PYTHONPATH=$CAFFE_ROOT/python
```

where *\$CAFFE\_ROOT* is the Caffe home directory, we managed to fix the issues.

**Web Camera.** Jetson TX1 on-board CSI camera does not work straightaway. On the other hand, plugging a dedicated web-camera almost does. To run Caffe SSD web-camera demo, add the following line before the command:

```
LD_PRELOAD=/usr/lib/aarch64-linux-gnu/  
↳ libv4l/v4l2convert.so
```

**Performance Tuning.** The new Tegra Linux driver package releases include *jetson\_clocks.sh* script, this is able to maximize performance by disabling DVFS, CPU idle, and CPU quit [10]. To toggle performance:

```
sudo ./jetson_clocks.sh
```

We recommend reading the manual first.

We also noticed that Jetson TX1 fan does not work on default. The script above turns it on. Enabling the fan without running the *jetson\_clocks.sh* script, can be achieved with:

